

# SupGen

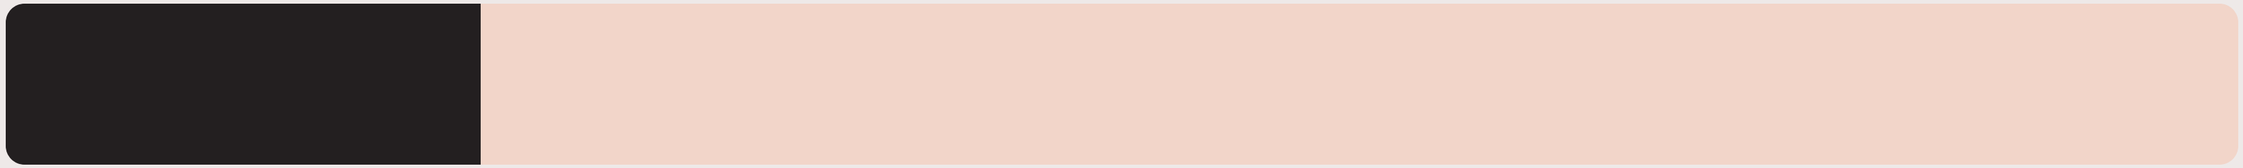
A different route to AGI

HIGHER ORDER COMPANY

PROBLEM

# AGI isn’t here yet

**The best AI today, on real paid freelance projects**



**21%** delivered at human level

**79%** not delivered

OpenAI defines AGI as *“highly autonomous systems that outperform humans at most economically valuable work.”*
Every frontier model is built on the transformer, and the transformer doesn’t generalize.

Sources: OpenAI Charter ↗ · Remote Labor Index, Scale AI & CAIS, top model as of Nov 2026 ↗ · Transformer: Attention Is All You Need, Vaswani et al., 2017 ↗

SOLUTION

# An AI architecture that generalizes

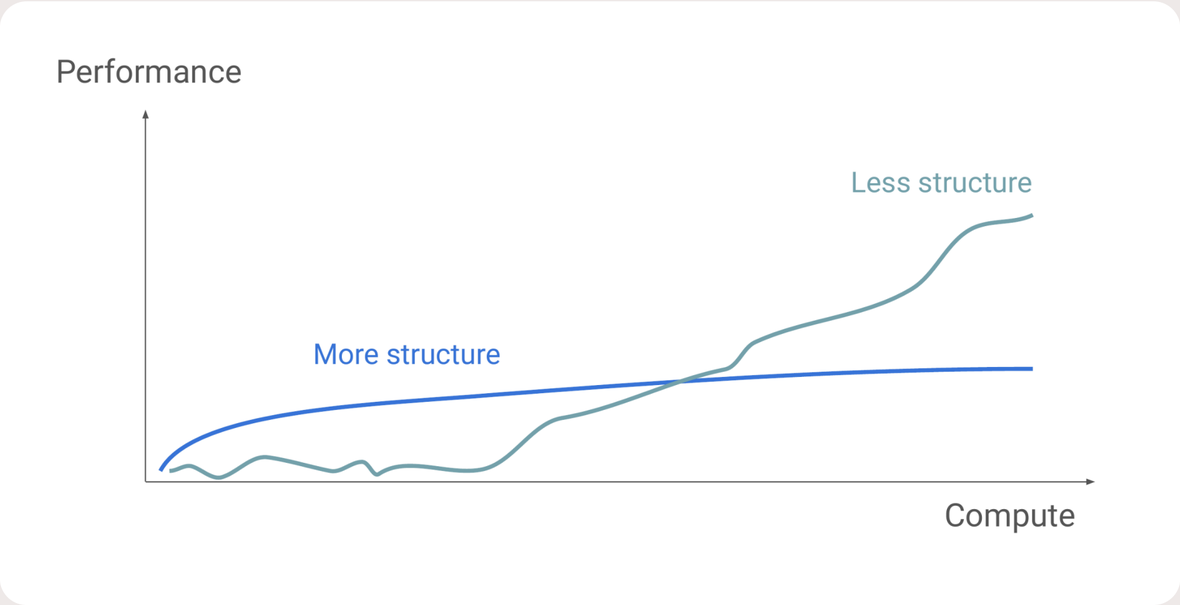



Figure: H. W. Chung, “Don’t teach. Incentivize.”, MIT EI Seminar, 2024, slide 9.

> *“The story of deep learning is this: empirically, old simple methods, which were usually invented in the 80s and the 90s, when scaled up on very large clusters, work really well. That happened with computer vision, that happened with normal supervised learning, and it also happened with RL.”*
>
> — **Ilya Sutskever**, co-founder of OpenAI and SSI, in a talk

## We’ll scale Symbolic AI to GPU clusters using HVM and Bend

BUSINESS MODEL

# We sell useful intelligence



**Elon Musk**
@elonmusk · Jun 18, 2026

Anthropic has rightly focused on maximizing useful intelligence, which does not show up in benchmarks, but definitely shows up in revenue.

**If research works, business is straightforward**

RAISING

# Up to $8M

On SAFEs, at a $120M post-money cap

**$100k** minimum check

**First come**, first served: investors@higherorderco.com

All the money goes to **research**

> This is a high-risk investment. You are investing in whatever comes out of our symbolic AI research; Bend will likely remain open source and free. Most early-stage investments return nothing, and we don’t yet have a clear long-term revenue path. We believe in our technology and our hustle, but that is no guarantee of success.

TEAM

# The team

**Victor Taelin**, founder, CEO, created HVM, Bend and Kind

**Vitor Chiarelli Neves**, cofounder

**Nicolas Abril**, researcher

**Lorenzo Battistela**, researcher

**Paulo Cavalcanti**, researcher

# Vision

The bitter lesson: the more humans hand-engineer an AI, the better it does at small scale, and the worse it scales. In the long run, the simplest method with the most compute wins. So we take that to the extreme. The simplest method there is: proof synthesis, run on massive compute.

That space holds every possible AI, including ones far better than anything a human could design, or even understand. The only question is how much compute it takes to find one. A learner like nanoGPT may fit in about 1,200 bits. Made by humans, it is surely far from the best. If the best is much smaller, an NVL72 can plausibly mine it.

It’s a big bet, on purpose. At $120M, success means about $2T: over 10,000x. Any chance above 0.01% makes it a good bet. We think it’s far higher, for the same reason GPT-2 went from a useless autocomplete to the heart of the AI industry: scale.

